



Pengui Labs

Building the next generation of AI-powered developer
infrastructure

Series A — 2026 San Francisco, CA

THE PROBLEM

AI development is broken

Modern teams juggle 8+ tools daily, losing 40% of engineering capacity to integration overhead and context switching.

- ✗ Fragmented tooling landscape
- ✗ No unified developer experience
- ✗ Manual orchestration of AI pipelines

40%

Engineering time wasted on infra plumbing

23

Context switches per engineer per day

8+

Tools in an average AI developer workflow

TRACTION

Key Metrics

6 months post-launch, Pengui Labs has established a strong growth trajectory across all key indicators.

\$2.4M

Annual Recurring Revenue
↑ 30% MoM

12K+

Active Developers
↑ 45% QoQ

72

Net Promoter Score
Industry avg: 34

99.97%

Platform Uptime

<50ms

Avg Response Time

3.2x

Developer Productivity Gain

Four pillars of Pengui



Pipeline Orchestration

Define complex AI pipelines in YAML. Automatic dependency resolution, parallel execution, and intelligent caching.

Core



Context Engine

Intelligent context management across your entire codebase. Semantic search, automatic chunking, and RAG-optimized retrieval.

Intelligence



One-Click Deploy

Ship AI-powered features in minutes, not weeks. Built-in scaling, versioning, A/B testing, and rollback capabilities.

Deploy



Observability

Real-time monitoring of model performance, cost tracking, latency analysis, and drift detection across all your AI systems.

Monitor



The future of AI development starts here

Join 12,000+ developers already building with
Pengui. Let's shape the future of intelligent
infrastructure together.

[Schedule a Meeting](#) [View Docs](#)